```
In [ 1:
       import gradio as gr
       import torch
       from transformers import AutoTokenizer, AutoModelForCausalLM
       # Load model and tokenizer
       model_name = "ibm-granite/granite-3.2-2b-instruct"
       tokenizer = AutoTokenizer.from_pretrained(model_name)
       model = AutoModelForCausalLM.from_pretrained(
           model name.
           torch_dtype=torch.float16 if torch.cuda.is_available() else torch.float32,
           device_map="auto" if torch.cuda.is_available() else None
       )
       if tokenizer.pad_token is None:
           tokenizer.pad_token = tokenizer.eos_token
       def generate_response(prompt, max_length=1024):
           inputs = tokenizer(prompt, return_tensors="pt", truncation=True, max_length=512)
           if torch.cuda.is available():
               inputs = {k: v.to(model.device) for k, v in inputs.items()}
           with torch.no_grad():
               outputs = model.generate(
                   **inputs,
                   max_length=max_length,
                   temperature=0.7,
                   do_sample=True,
                   pad_token_id=tokenizer.eos_token_id
           response = tokenizer.decode(outputs[0], skip_special_tokens=True)
           response = response.replace(prompt, "").strip()
           return response
       def disease_prediction(symptoms):
           prompt = f"Based on the following symptoms, provide possible medical conditions and general medication suggestions. Alw
           return generate_response(prompt, max_length=1200)
       def treatment_plan(condition, age, gender, medical_history):
           prompt = f"Generate personalized treatment suggestions for the following patient information. Include home remedies and
           return generate_response(prompt, max_length=1200)
       # Create Gradio interface
       with gr.Blocks() as app:
           gr.Markdown("# Medical AI Assistant")
           gr.Markdown("**Disclaimer: This is for informational purposes only. Always consult healthcare professionals for medical
           with gr.Tabs():
               with gr.TabItem("Disease Prediction"):
                   with gr.Row():
                       with gr.Column():
                           symptoms_input = gr.Textbox(
                               label="Enter Symptoms",
                               \verb|placeholder="e.g., fever, headache, cough, fatigue...",\\
                           predict_btn = gr.Button("Analyze Symptoms")
                       with gr.Column():
                           prediction_output = gr.Textbox(label="Possible Conditions & Recommendations", lines=20)
                   predict_btn.click(disease_prediction, inputs=symptoms_input, outputs=prediction_output)
               with gr.TabItem("Treatment Plans"):
                   with gr.Row():
                       with gr.Column():
                           condition_input = gr.Textbox(
                               label="Medical Condition"
                               \verb|placeholder="e.g., diabetes|, hypertension, migraine...",\\
                               lines=2
                           )
                           age_input = gr.Number(label="Age", value=30)
                           gender_input = gr.Dropdown(
                               choices=["Male", "Female", "Other"],
                               label="Gender",
                               value="Male"
                           history_input = gr.Textbox(
                               label="Medical History",
                               placeholder="Previous conditions, allergies, medications or None",
                               lines=3
                           )
```

```
plan_btn = gr.Button("Generate Treatment Plan")
                              with gr.Column():
                                     plan_output = gr.Textbox(label="Personalized Treatment Plan", lines=20)
                       plan_btn.click(treatment_plan, inputs=[condition_input, age_input, gender_input, history_input], outputs=plan_c
 app.launch(share=True)
 /usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your
 Google Colab and restart your session.

You will be able to reuse this secret in all of your notebooks.

Please note that authentication is recommended but still optional to access public models or datasets.
    warnings.warn(
tokenizer_config.json: 0.00B [00:00, ?B/s] vocab.json: 0.00B [00:00, ?B/s] merges.txt: 0.00B [00:00, ?B/s]
tokenizer.json: 0.00B [00:00, ?B/s] added_tokens.json: 0%|
added_tokens.json: 0%| | 0.00/87.0 [00:00<7, ?B/s]
special_tokens_map.json: 0%| | 0.00/701 [00:00<7, ?B/s]
config.json: 0%| | 0.00/786 [00:00<7, ?B/s]
`torch_dtype` is deprecated! Use `dtype` instead!
```

 model.safetensors.index.json:
 0.00B
 [00:00, ?B/s]

 Fetching 2 files:
 0%|
 0/2 [00:00<?, ?it/s]</td>

 model-00001-of-00002.safetensors:
 0%|
 0.00/5.1M [00:00<?, ?B/s]</td>

 model-00002-of-00002.safetensors:
 0%|
 0.00/67.1M [00:00<?, ?B/s]</td>

 Loading checkpoint shards:
 0%|
 0/2 [00:00<?, ?it/s]</td>

 generation_config.json:
 0%|
 0.00/137 [00:00<?, ?B/s]</td>

Colab notebook detected. To show errors in colab notebook, set debug=True in launch() * Running on public URL: https://88743294eab45262a0.gradio.live

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working directory to deploy to Hugging



No interface is running right now